


False human premise in the first prompt



Selected questions previously answered correctly by both models. Error = wrong final answer; no explicit abstentions.

Each model: 50 independent neutral / false-premise prompt pairs. Bars are descriptive; exact discordance $p = 0.25$ (DeepSeek), 0.125 (MiniMax).